



Data Science for Linguists

Session 4: Data Cleaning and Preparation

Johannes Dellert

17 November, 2023



Data Cleaning and Preparation

- in data science, data preparation (cleaning, transforming, and rearranging of real-world data) often comprises the majority of the work; a common figure is about 80% of an analyst's time
- ad hoc processing of data using a variety of tools (plain Python, Perl, R, Java, sed, awk) is still very common in science, but this approach has many disadvantages
- using Pandas in a Jupyter notebook for these tasks has the advantages of
 - ▷ a high-level, flexible, and fast set of tools that is very accessible to other researchers
 - ▷ full documentation of the steps taken to derive the research data from the raw inputs (crucial for transparency and reproducibility!)



Table of Contents

Handling Missing Data

Duplicate Removal, Value Replacement, and Renaming

Discretisation and Binning

Outlier Detection and Removal

Permutation and Random Sampling

Vectorised String Operations

Assignment 4



NaN and None in Pandas

- due to their underlying implementation in NumPy, Pandas series represent missing values by means of different **sentinel values**, the choice of which depends on the datatype:
 - ▷ for data with dtype `float64`, Pandas uses the floating-point value NaN (Not a Number) which is defined by the ISO standard (accessible via `np.nan`)
 - ▷ for general Python objects (including variable-length strings), it uses the special value `None` from standard Python (which is very slow to compute with!)
 - ▷ because these two options are the only ones available, series of integer datatypes are silently converted into floating point numbers as soon as missing values are inserted
- one of the goals of Pandas is to make working with missing data efficient and smooth
 - ▷ there is well-defined default behaviour for any functions executed on Pandas series with missing values; such values will never cause Pandas to throw an exception, though the default behaviour might not be extremely useful in some cases
 - ▷ many methods (such as the ones for descriptive statistics) are implemented to silently exclude missing data per default
 - ▷ it also provides a range of methods which abstract over the somewhat inconsistent underlying implementation of null values



Pandas Nullable Dtypes

- to add support for true integer arrays with missing data, Pandas provides **nullable dtypes** like `pd.Int32` (distinguished from the default dtypes like `pd.int32` by capitalisation)
- the null value in series of these types is represented by `pd.NA`
- the other null values will be normalised without triggering implicit typecasting:

```
In [ ]: pd.Series([1, np.nan, 2, None, pd.NA], dtype='Int32')
```

```
Out[ ]: 0      1
```

```
       1  <NA>
```

```
       2      2
```

```
       3  <NA>
```

```
       4  <NA>
```

```
dtype: Int32
```

- for more efficient handling of large amounts of string data, there is the specialised extension type `pd.StringDtype()`



Detecting Null Values

- null values of any kind can be detected using the method `data.isnull()`, which returns a Boolean mask over the data, with `True` in all positions where there is a null value
- `data.notnull()` returns a Boolean mask as well, but the opposite of the result of the previous method (`False` in all positions where there is a null value)
- `data.isna()` and `data.notna()` are aliases of `data.isnull()` and `data.notnull()`, both sentinel values are matched by both methods (unlike in the underlying implementation)

```
In [ ]: data = pd.Series([1, np.nan, 'hello', None])
In [ ]: data.isnull()
Out[ ]: 0    False
        1     True
        2    False
        3     True
        dtype: bool
```



Dropping Null Values with `dropna()`

- `data.dropna()` is a convenience method which combines filtering for and then dropping the null values, i.e. a shorthand for `data[data.notna()]`
- if `data` is a `Series`, the result will simply be a shorter series
- if `data` is a `DataFrame`, we can only drop entire rows or columns
 - ▷ by default, `data.dropna()` will drop all rows which contain any null value
 - ▷ to drop columns instead, we can provide the argument `axis=1` or `axis="columns"`
- argument `how="all"` to only drop rows/columns which consist entirely of null values
- for more fine-grained control (e.g. only weeding out the most gappy records), we can also provide an argument `thresh=k` to specify that any row/column with at least k non-null values will be kept



Filling Null Values with `fillna()`

- `data.fillna(x)` is a convenience method which combines filtering for and then overwriting the null values with a default value `x`, i.e. a shorthand for `data[data.isna()] = x`
- `x` can be a dictionary providing different fill values by column index
- `data.fillna(method="ffill")` specifies a forward fill, i.e. empty values will be replaced by the previous non-null value (which therefore gets propagated forward)
 - ▷ example: a series with data `[1 <NA> 2 <NA> <NA> 3]` will become `[1 1 2 2 2 3]`
- `data.fillna(method="bfill")` specifies a backward fill, i.e. empty values will be replaced by the subsequent non-null value (which therefore gets propagated backward)
 - ▷ example: a series with data `[1 <NA> 2 <NA> <NA> 3]` will become `[1 2 2 3 3 3]`
- if no previous or subsequent value is available, `ffill` and `bfill` leave null values at the fringes!
- for a `DataFrame`, we can again switch to propagation through columns by specifying `axis=1`



Table of Contents

Handling Missing Data

Duplicate Removal, Value Replacement, and Renaming

Discretisation and Binning

Outlier Detection and Removal

Permutation and Random Sampling

Vectorised String Operations

Assignment 4



Removing Duplicates

- `data.duplicated()` returns a Boolean Series indicating whether each row is a duplicate (all values equal to some previous row) or not
- `data.drop_duplicates()` returns a Series or DataFrame consisting of only those rows in data where `data.duplicated()` was False
- the `subset` argument allows to provide a list of column indices to specify which columns are relevant for duplicate detection (other columns are allowed to have different values)
- if the `subset` argument is used, the first variant of each duplicate will be used for the values in the irrelevant columns, `keep="last"` changes this



Replacing Values

- `data.replace(oldvals, newvals)` substitutes a set of values by replacements
 - ▷ `data.replace(-1, 0)` sets all cells with value -1 to 0
 - ▷ `data.replace([-2, -1], 0)` sets all cells with value -2 or -1 to 0
 - ▷ `data.replace([-2, -1], [-1, 0])` sets cells with value -2 or -1, and all with -1 to 0
- `data.replace(dictionary)` replaces each occurrence of a key with its value
- more complex element-wise transformation can be implemented as a function (e.g. `transform_value(x)`), and then executed on every cell by a call to `data.map(transform_value)`; this works with anonymous functions (`lambda`) as well



Renaming Axis Indexes

- the axes can be modified in place by executing the map method of their indices:
 - ▷ `data.columns = data.columns.map(str.title)`
 - ▷ `data.index = data.index.map(lambda x: x[:4].upper())`
- `data.rename()` allows to create a transformed version of a dataset without modifying the original (simple example: `data.rename(index=str.title, columns=str.upper)`)



Table of Contents

Handling Missing Data

Duplicate Removal, Value Replacement, and Renaming

Discretisation and Binning

Outlier Detection and Removal

Permutation and Random Sampling

Vectorised String Operations

Assignment 4



Discretisation and Binning

- continuous data is often discretised or otherwise separated into bins for analysis
- `pd.cut(data, k)` returns a `Categorical` object which describes the k equal-length bins computed on the basis of the minimum and maximum values within the data
- `pd.cut(data, cutoffs)` returns a `Categorical` object which describes the bins computed from the data based on the specified cutoff values between the bins
- `pd.qcut(data, k)` bins the data into k quantiles (equally sized bins)
- `pd.qcut(data, quantiles)` bins the data into the provided quantiles (between 0 and 1)
- the argument `labels=group_names` allows to override the default bin names
- a `Categorical` object `cats` has the following key applications:
 - ▷ `cats.codes` returns an array containing the bin index for each datapoint
 - ▷ `cats.categories` shows an `IntervalIndex` object representing the bins
 - ▷ `pd.value_counts(cats)` renders the counts of datapoints in each bin
- these functions will have central importance in Session 7 (aggregation and grouping)



Table of Contents

Handling Missing Data

Duplicate Removal, Value Replacement, and Renaming

Discretisation and Binning

Outlier Detection and Removal

Permutation and Random Sampling

Vectorised String Operations

Assignment 4



Detecting and Filtering Outliers

- outlier detection and filtering is typically performed by combinations of simple array operations
- for a normally distributed Series of data, we might define our outliers as all rows where the value exceeds 3 in absolute value: `data[data.abs() > 3]`
- in a DataFrame full of normally distributed values, we might be interested in the rows where any of the columns has such a value: `data[(data.abs() > 3).any(axis="columns")]`
(Boolean DataFrame generated by the comparison, on which we apply the `any()` method)
- to remove the outliers by capping values to the range `[-3, 3]`, we can just do `data[data.abs() > 3] = np.sign(data) * 3`



Table of Contents

Handling Missing Data

Duplicate Removal, Value Replacement, and Renaming

Discretisation and Binning

Outlier Detection and Removal

Permutation and Random Sampling

Vectorised String Operations

Assignment 4



Permutation and Random Sampling

- permutations are sampled using `smp = np.random.permutation(k)`
- `data.take(smp)` (= `iloc`-based indexing) is then a random permutation of the first k lines
- `data.sample(n=k)` selects a random subset of k rows without replacement
- `data.sample(n=k, replace=True)` samples k rows with replacement



Table of Contents

Handling Missing Data

Duplicate Removal, Value Replacement, and Renaming

Discretisation and Binning

Outlier Detection and Removal

Permutation and Random Sampling

Vectorised String Operations

Assignment 4



Vectorised String Operations: Equivalents of Basic Methods

- cleaning up a messy dataset often requires a lot of string manipulation, but simple element-wise application using `data.map()` will fail on the null values
- the `Pandas Series` offers array-oriented and null-aware string operations which are accessible via the `str` attribute; here is a small sample:
 - ▷ `data.str.count` to count occurrences of a pattern
 - ▷ `data.str.contains(s)` returns a Boolean mask with the result of calls to `s in value` on each cell value, mixed with NaN values wherever a value was missing
 - ▷ `data.str.len` to compute the length of each string
 - ▷ `data.str.strip` to trim whitespace (including newlines) from both sides



Vectorised String Operations: Regular Expressions

- there are also vectorised and null-aware versions of the regex capabilities:
 - ▷ `matches = data.str.findall(r"somePattern")`
 - ▷ `first_matches = matches.str.get(1)`
- `data.str.extract(pattern)` returns the captured groups of the regular expression as a new DataFrame



Vectorised String Operations: Miscellaneous Methods

- several methods emulate the capabilities of Python string operations:
 - ▷ `data.str.cat` for element-wise concatenation (with optional delimiter)
 - ▷ `data.str.get` to index each string
 - ▷ `data.str.repeat` to repeat each string
 - ▷ `data.str.slice` for extracting slices from each string
- `data.str.get` and `data.str.slice` are also available through the normal indexing syntax (`data.str[i]` and `data.str[i:j]`)



Table of Contents

Handling Missing Data

Duplicate Removal, Value Replacement, and Renaming

Discretisation and Binning

Outlier Detection and Removal

Permutation and Random Sampling

Vectorised String Operations

Assignment 4



Assignment 4: Tasks

- 1) Read up on the CoNLL-U format for dependency treebanks.
- 2) Create a new Jupyter notebook and import pandas. Load the development set of the Basque UD corpus from the file `eu_bdt-ud-dev.conllu` into a DataFrame object, filtering out lines which are empty or start with `#`. Specify the CoNLL-U field names as a column index.
- 3) Convert all values consisting of an underscore into an appropriate missing data type, and remove all columns where more than 80% of the values are empty.
- 4) Reduce the dataset to all rows representing forms of the auxiliary *izan* (“to be”).
- 5) Apply vectorised string operations which involve regular expressions to convert the morphological features into a more useful format (one new column per feature, feature values in each row; example: column `Number[abs]` with values `Sing` and `Plur`).
- 6) Query the database to see whether there is any syncretism in the paradigm of *izan*. (i.e. same form, different morphological features) [Hint: duplicate detection on different subsets]
- 7) How often does each form of *izan* occur in this development set? Create a new dataframe consisting of the form, the columns with the morphological features, and these counts.
- 8) Split the forms of *izan* attested in the corpus into ten bins of equal size, so that each bin groups together forms of roughly equal frequency. Detect and remove outliers if necessary.



Preliminary Course Plan

- 1 27/10 IPython and Jupyter**
- 2 03/11 Introduction to NumPy**
- 3 10/11 Pandas and Data Frames**
- 4 17/11 Data Cleaning and Preparation**
- 5 24/11 Linguistic Preprocessing
- 6 01/12 Data Wrangling: Join, Combine, Reshape
- 7 08/12 Data Aggregation and Grouping
- 8 15/12 Visualisation with Seaborn
- 9 22/12 Modeling and Prediction
- 10 12/01 Classification
- 11 19/01 Clustering
- 12 26/01 Pattern Extraction and Density Estimation
- 13 02/02 Statistical Inference
- 14 09/02 Data Science Projects



Questions

Questions?

Comments?

Suggestions?